litellm vs agentgateway

Run 20260626-120716 · 32 connections · 3s duration · max QPS · 1,104-byte payload · 18 LiteLLM workers · mock OpenAI backend

Source: fortio 1.73.0 · docker stats · results/20260626-120716/

ⓘ agentgateway handled 11.5× more requests per second with 8.5× lower median latency, while using ~540× less memory. Both gateways returned HTTP 200 for every request.

36,934 req/s
Throughput (agentgateway)

3,198 req/s
Throughput (litellm)

22.4 MiB
Avg memory (agentgateway)

11.8 GiB
Avg memory (litellm)

Throughput and latency

Throughput comparison req/s

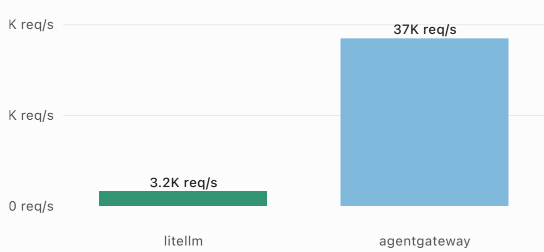

Y-axis: requests per second (req/s). X-axis: gateway under test.

Latency percentiles ms

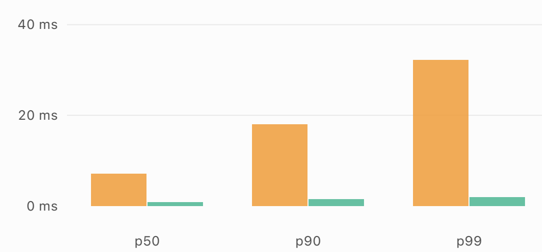

Y-axis: end-to-end latency (ms). Lower is better.

CPU and memory

Sampled every 1s via docker stats during the 9s benchmark window (4 samples).

CPU usage under load % of one core

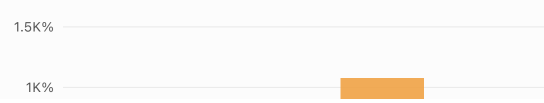

Memory usage under load MiB

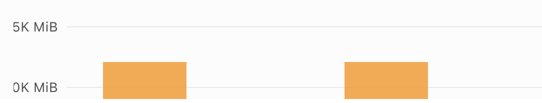

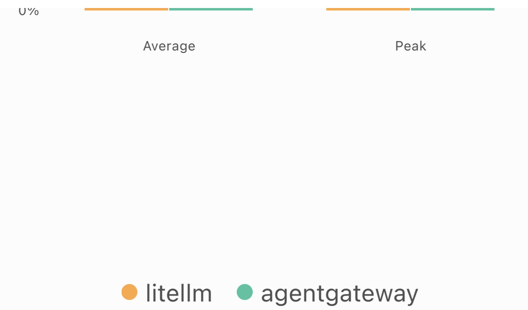

Y-axis: CPU percent (can exceed 100% with multiple cores).

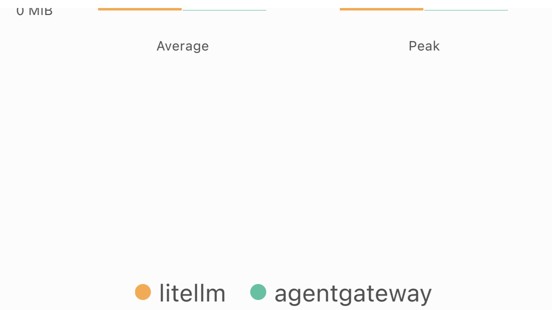

Y-axis: resident memory (MiB). litellm uses ~11.8 GiB with 18 workers.

Latency distributions

Request-count histogram buckets from fortio. agentgateway clusters sub-millisecond; litellm spans single-digit to tens of milliseconds.

litellm latency histogram 10,186 req

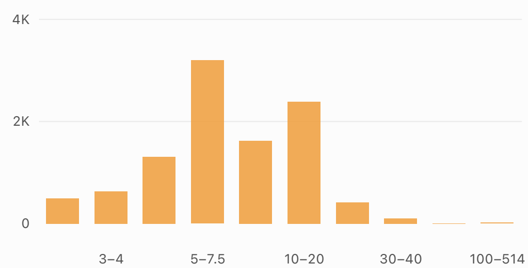

X-axis: latency bucket (ms). Y-axis: request count.

agentgateway latency histogram 110,831 req

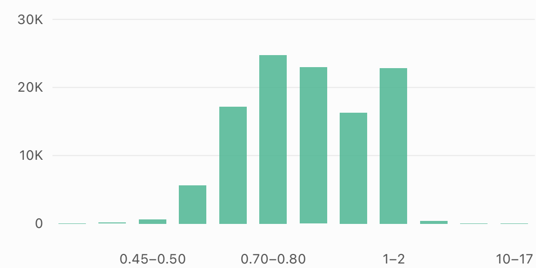

X-axis: latency bucket (ms). Y-axis: request count.

Latency summary

Metric	litellm	agentgateway
Throughput	3,198.48 req/s	36,933.62 req/s
Total requests (3s)	10,186	110,831
p50 latency	7.08 ms	0.83 ms

CPU and memory summary

Metric	litellm	agentgateway
Avg CPU	330.79%	104.80%
Peak CPU	1075.38%	243.02%
Avg memory	11.81 GiB	22.38 MiB
Peak	11.81 GiB	28.79 MiB